

BACHELOR INFORMATICA

UvA  UNIVERSITEIT VAN AMSTERDAM

Dynamic program loading in a shared address space

Leendert van Duijn

May 21, 2012

Supervisor(s): Raphael 'Kena' Poss (UvA)

Signed: Raphael 'Kena' Poss (UvA)

Abstract

In this abstract the research will be summarized to some extent.

Contents

1	Introduction	4
2	Theoretical background	5
2.1	The platform	5
2.1.1	Features	5
2.2	The loading	5
2.2.1	User input	5
2.2.2	Loading from ELF	6
2.2.3	Location decisions	6
2.2.4	Relocation	6
2.2.5	Process private memory	6
2.2.6	Execution	6
2.3	User control	7
2.3.1	Preparation and Compilation	7
2.3.2	Configuration details	7
3	Implementation	8
3.1	Assumptions and constraints	8
3.1.1	C compiler	8
3.1.2	The linker	8
3.1.3	Flags for linker and compiler	8
3.2	Api	9
3.3	Platform dependency	9
3.4	Configuration	9
3.5	Elf loading	9
3.5.1	Special symbols	9
3.5.2	Algorithm	10
3.5.3	Program limitations and requirements	11
3.6	Spawning an initial program	12
3.7	In program Loader calls	12
3.8	Implementation considerations and reflection	12
3.8.1	Location dilemma	12
3.8.2	Trace ability of problems	12
3.8.3	Relocation	12
4	Progress report	13
4.1	Milestones	13
4.1.1	Planned milestones	13
4.1.2	Unexpected roadblocks	13
4.1.3	Reached milestones	13
4.2	Future research	14
4.3	Security	14

5	Experiments	15
5.1	Testing	15
5.1.1	Relocation	15
5.2	Limitations and future work	15
5.2.1	Permissions	15
5.3	Applications	15
6	Conclusions	16
6.1	Benchmarking results	16
6.2	Stability	16
6.3	Final conclusion	16
6.4	References	16
A	Problems and in depth solutions	18
A.1	Bugs	18
A.1.1	String data relocation	18
A.2	Example Configurations	19

Introduction

When a computer loads a program for execution it is granted its own private address space. In this research we hope to show that more than just a single program can cooperate in a shared address space. This opens up possibilities for processes designers to share hardware between an increasingly larger number of cores without limiting the number of distinct programs.

Theoretical background

2.1 The platform

Microgrid. A simulation platform for a massively parallel computation platform consisting of multiple general purpose processors based on the DEC Alpha processor architecture. This platform is designed to explore the possibilities of increasingly parallel computing systems focusing on a set ¹ of general purpose processing units capable of cooperation at very low cost and latency. This is hoped to increase trough put...

2.1.1 Features

The Alpha processor

- 64bits address space
- General purpose CPU design
- Risc instruction set

The Microgrid environment has several key features which have led to this research specifically.

- Parallel computation
- Large address space and memory pool
- TODO share a TLB
- Houses more parallel processing power than a typical application needs...

2.2 The loading

Loading a program has several phases to go through. At the end of these steps a traditional loader would transfer full control to the loaded software.

2.2.1 User input

A user needs to tell the loader what programs need to be loaded and what specific parameters or settings should be used. This can done by means of a simple configuration file as detailed in Section 2.3.2.

¹Documentation about: How many procs in microgrid

2.2.2 Loading from ELF

The loading of an ELF executable as specified in [1]. Loading would result in a set of memory ranges being populated with code and data. Administrative features included in the ELF format like the program entry point and symbol data are used beyond this point though they are not necessarily part of a fully loaded program.

2.2.3 Location decisions

Since programs will share the address space in a parallel fashion a single arbiter needs to decide where a process can be loaded. In order to retain high performance with an increasing amount of processes a freelist could be maintained. This list points to an administrative entry which is guaranteed to be either available or a truthful indicator that there is no space whatsoever. It contains a link to the next available entry. On process termination, its entry can be attached to the front of the freelist ensuring that all memory ranges are accounted for at any time.

2.2.4 Relocation

During the loading process an exact location is determined for the program. This location is highly dynamic. When the program is requested to load any other program load can affect its final location. This introduces the need for program relocation ². The relocation can be split into two important phases. The code relocation and data relocation. The code is not always trivially relocated ³. To protect the scope of this research the loader demands for the loaded programs to be compiled with several flags related to Position Independent Codeso that the code is functionally independent of its location in memory. The needed flags are elaborated on in Section 3.5.3. This leaves some data relocation entries to be processed by the loader in order to correct data pointers⁴ These relocation corrections are done as specified in ⁵. This is summarized as adding the programs base address to each pointer the compiler has flagged for correction. These pointers are full size pointers which places no extra limitations on program location. Some of these pointers could be function pointers for usage by the Position Independent Codethis is however irrelevant to the relocation code.

2.2.5 Process private memory

As programs may require arguments and environment variables which outlast the parent process. This requires storage allocation. This allocation can be supported by the ELF file format by including a special section which reserves space for either arguments, environment or other custom data segments. This section is detected when loading and if present will be used to pass any argument and environment variables. If this section is absent no arguments will be passed to the program. This enables a minor speedup and memory saving for programs which are known not to use arguments⁶.

2.2.6 Execution

Transferring control to the loaded program. This is done via the Microgrid function `sl_create` which places the program on the desired cores, taking as parameters the address to start execution at and any arguments to pass. It offers an option to fully reserve the cores for the given program, blocking any core sharing. Due to the serializing effects of this option and the absence of truly intelligent core selection this is disabled by default.

²Documentation about: position dependent code

³Documentation about: coderelocationhard

⁴For an example and explanation see Section A.1.1.

⁵Documentation about: elf reloc data reference

⁶Documentation about: example no arg programs and speedup

2.3 User control

The loader can be influenced by a user in several ways. During its preparation and compilation several settings can be tweaked for optimum performance as will be specified in Section 2.3.1. After the loader has been compiled and linked the remaining tweaks need to be done via configuration, most tweaks are program specific so that once a program is loaded other ill-written programs can not legally affect it⁷.

2.3.1 Preparation and Compilation

During compilation some values can be tweaked. Such as the memory size for programs, the default cores for some operations, debugging prints.

- Ranges Base
- Ranges Size
- Maximum number of programs
- Core used for serialized printing
- Core used for administrative functions

2.3.2 Configuration details

The user can guide the loader by using a quite powerful configuration file. The file can be easily written in any text editor capable of saving as plain formatted text. Some obligatory settings

- Filename of the ELF file (string)
- Arguments, can be an empty line (newline separated strings)
- Environment, can be an empty line (newline separated strings)

Some settings are optional but would be present in most cases

- Verbose, "true" or a numerical value"
- Exclusive, "true" or "false" (Optional, default false, TODO)
- Core_start, numerical core number, (Optional)
- Core_size, numerical number of cores, (Optional, defaults to 1)

⁷Some efforts of sabotage are predicted to be effective as detailed in ⁸

Implementation

3.1 Assumptions and constraints

In the implementation process of the loader several assumptions had to be made, most of which will be demonstrated to be correct.

- Availability of a working C compiler, gcc
- Availability of a working linker, gcc
- Availability of the -fpic -fPIC and -shared flags
- Correctness of loaded code
- Relocateability of loaded code
- Loaded code will not try to harm other loaded programs

3.1.1 C compiler

The C compiler used during development is the SLC ¹compiler which has shown to compile an extensive set of test programs such as the benchmark set. . .

3.1.2 The linker

In order to link the loader, which is divided over a set of source and header files. SLC is used which is fully compatible with the used C compiler. This linker is primarily tasked with generating object files from the loader source and linking these together in a self contained executable. Optionally including relocation information if runtime relocation is desired.

3.1.3 Flags for linker and compiler

The C compiler and linker need to follow some specific rules in order to maintain full relocateability and functional correctness. These flags are the -fpic -fPIC and -shared flags. The -f flags indicate to the compiler that any executable code needs to be fully relocateable. The -shared flags indicated to the linker all data references might be relocated prior to execution and as such administration to support should be included. These flags are needed to compile any program that should reliably run within our loader.

To implement the loader the C programming language is employed, a compiler converts this into executable code for the Microgrid platform.

¹Documentation about: slc 3.7b.28-dac6

Copyright (C) 2009,2010 The SL project. This is free software; see the source for copying conditions. There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

Written by Raphael 'kena' Poss.

3.2 Api

The loader passes a pointer to all loaded programs which can be used to access loader functions. This pointer points to a struct containing function pointers for these functions:

- `spawn`, spawn a program
- `print_string`, prints in an orderly fashion (blocking)
- `print_int`, prints an integer in an orderly fashion (blocking)
- `print_pointer`, prints a pointer in an orderly fashion (blocking)
- `load_fromconf`, loads a program from config file
- `load_load_fromconf_fd`, loads a program from config from file descriptor
- `load_fromparam`, loads from structure with parameters
- `breakpoint`, TODO loader break point for single program

3.3 Platform dependency

We depend on the microgrid for several key functions, these would need to be replaced if any other platform were to be targeted.

- `sl_create`, for creating the program stack and core allocation
- `sl_detach`, letting the loader detach from a loaded program
- `mgsim_control`, for memory range allocation and deallocation

3.4 Configuration

A simple scanner which parses keyword value pairs, these are then terminated by a blank line at which point the arguments can be specified. These arguments will be passed as the traditionally called `argv`, which can only be done if the necessary room is reserved in a section as detailed in Section 3.5.1. These newline separated arguments are terminated by a blank line. After this blank line the environment variables are once written separated by newlines. The environment variables should be in the form `a=b`. The environment variables are terminated by a blank line after which any remaining data would be left untouched.

3.5 Elf loading

3.5.1 Special symbols

The loader searches for some special symbols which it can use to store and pass arguments to programs in an unobtrusive manner. These symbols are generated by including a C source file during compilation² of the loadable program. These are symbols with global scope, which is global to the compiled program. Other loaded programs do not see them. These symbols are detected when parsing the dynamic symbol table and include both the size and the unrelocated location, after correcting for relocation the symbol location is stored in the programs administration for later use. The symbols are recognized by their names, these can be changed by altering the definition of `ROOM_ARGV` or `ROOM_ENV` in the file `loader_api.h` and the related C source file which would be either `argroom.c` or `envroom.c`. The latter also permit the size to be modified in order to accommodate for the anticipated amount of arguments.

²linking an object file compiled from this file will achieve the same effect

Size constraints and guidelines

The size the argument and environment objects require depends on the anticipated input, in order to calculate the most efficient size these formula should be used:

$$Size_{Env} = 1 + \sum_{i \in environment} (1 + strlen(i))$$

$$Size_{Args} = 8 * (Argc + 1) + \sum_{i \in argv} (1 + strlen(i))$$

The room needed for the arguments considers the storage for the argv array, the environment room does so for the final null byte. These storage locations are only related in concept and implementation. They are fully independent so one may choose to include any combination of sizes.

In the situation insufficient room is available for the arguments the loader will print a warning message, setup to pass no arguments whatsoever. It will then check the same for the environment variables. It will still try to execute the program even if these checks both fail by passing null pointers and an argc of zero to indicate no arguments could be passed. It is left up to the developer of the loaded program to decide whether it can successfully execute in their absence.

3.5.2 Algorithm

The Elf file is read into memory where a simple algorithm is followed.

```
Load the file into memory
Inspect the header
Locate the program headers
Scan the program headers for the base address
Find an available PID
Determine the read base address, aligned
Loop over the program headers:
    Load segments to their destination
    Zero leftover memory
Locate the section headers
Scan the section headers for the Dynamic Symbol table and Relocation tables
Scan the Dynamic Symbol table for special symbols:
    Note the location and size for the argument and environment room
For all found Relocation tables:
    Loop over all entries in the table:
        Determine the pointer location
        Determine the symbol and offset
        Calculate and update the pointer
Prepare the arguments and environment if room is available
Spawn a thread which will call the main function as found in the entry point
```

The loader optionally includes a verbose set of print statements useful for debugging purposes. This can be disabled for performance reasons by changing a macro definition or disabled at runtime by passing a verbosity setting to the loader.

The loader is guided by a configuration file which describes what program should be called with optional arguments and program specific settings. This configuration file is covered in detail in Section 2.3.2. The loader follows two simple algorithms for parsing.

Reading key value pairs:

Start:

Key=""

Value=""

Buffer=""

Read character X:

```

If X == '=':
    Key=Buffer
If X == '\n':
    if Key == "":
        Goto Done
    Value=Buffer
    Goto ParseSetting(Key, Value)
Buffer += X
Goto Read character X

ParseSetting(Key, Value):
    Pick the setting based on Key, set it using Value
    Return

Done:
    Finish up, settings done

```

At this point all settings have been parsed, the programs filename is known and the settings have been terminated with an empty line. At this point the command line arguments can be set..

```

Argc=1
Argv[0]=ElfFilename
Read character X:
    if X == '\n':
        if Argv[Argc][0] == '\0':
            Goto Done
        Argc++
    Argv[Argc] += X
    Goto Read character X
Done:
    Finish up, Arguments known in Argv and Argc

```

The same is the done for the Environment substituting Argc and Argv for EnvC and Evnp.

3.5.3 Program limitations and requirements

Some compilation flags and settings are explicitly required in order to reliably load a program.

- -fPIC
- -fpic
- -shared
- crt_fun.c
- -nostdlib

These flags tell SLC ³to compile position independent code, to include data relocation information and replace the C runtime codewith a bare one which consists of a wrapper which calls the main function of the loadable program.

³Documentation about: slc 3.7b.28-dac6

Copyright (C) 2009,2010 The SL project. This is free software; see the source for copying conditions. There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

Written by Raphael 'kena' Poss.

3.6 Spawning an initial program

The loader initial program is loaded by passing arguments which will be parsed as configuration files, loading them in sequence on either the default core or the specified cores. These files should adhere to the format as specified in Section 2.3.2. Several examples are included in Section A.2

3.7 In program Loader calls

In order to allow more complex program structure we offer programs an API through which they can invoke loader functions to spawn programs or perform other related tasks.

3.8 Implementation considerations and reflection

A perfect design is rare, as not all optimum settings are fully compatible.

3.8.1 Location dilemma

Size

A program needs a location which can not be easily changed at run time. The run time of a program may be unbounded and as such a single program introduces an obstacle in the memory space. This is largely an allocation problem where prior to execution exact space requirements may not be available.

Location

When loading a program care has to be taken to ensure no programs are given overlapping address spaces⁴.

During the debugging of any collection of programs one would like to know which program is responsible for certain instructions, problems or memory usage. In order to trace a specific memory location to a program we would need some sort of standard procedure.

Solution

As our loader is designed with a 64bits address space in mind we have adopted a formula for base address calculation in which a process is given a base address based on its identifier and a predetermined size. This size is the upper limit for any loaded process sub address space. This size can be changed prior to loader compilation.

$$Base = Base_{Global} + (Id_{Process} * Size_{maximumsubspace})$$

This enables us to efficiently determine a base address for a process and also pinpoint the source of many memory related problems as an offending instruction can be trivially traced back to its program based on its address.

3.8.2 Trace ability of problems

As our loader loads an increasingly large number of programs problems memory ownership needs to stay intuitive. In order to trace errors we can determine ownership by calculating the Id of the memory based on the inverse of our location formula.

3.8.3 Relocation

The loaded program is loaded to an a priory unknown location so the loader has to finish the relocation process. This is done by using relocation information information stored in the section headers. The loader parses the section data. [2]

⁴All programs share the same address space, but to their knowledge the subspace they inhabit is a traditional address space

Progress report

During my research I have reached several conclusions.

4.1 Milestones

4.1.1 Planned milestones

A loader for single program.
A loader for multiple programs.
A loader with in program spawn function.
A loader with Input and Output redirection.

4.1.2 Unexpected roadblocks

Missing functions.
Libc conflicts, loading an existing libc leads to crashes.
Runtime cleanup, programs not being cleaned due to thread termination.

4.1.3 Reached milestones

Loading a program.
Loading more programs.
Loading based on user configuration.
Loading on a specific core.
Letting programs output in a orderly fashion.

Loading something

A single program being loaded, though it seems trivial it is quite the relief when it finally does.

Loading several programs

The more significant milestone, loading multiple programs which execute as they would in a private address space. With the significant difference that they share their address space between themselves and the loader.

4.2 Future research

4.3 Security

As the loader is focused on sharing an address space between programs some assumptions were made as seen in ¹. One of these assumptions is the willingness

... What has been found to be unanswered for now... What could be better...

¹Documentation about: `assumtionsprograms`

Experiments

5.1 Testing

During the development several programs were written to test nominal behaviour. These programs are designed to make use of several features of the ELF file format which could break on loader malfunction.

5.1.1 Relocation

The `tinyex.c` prints strings which are globally defined in an array. This array of string pointers requires runtime relocation to ensure they point to the relocated string data. These are full size absolute pointers and as such corrected by simple addition of the program base offset which is unknown at compile time.

Symptoms of malfunction for this program would include illegal memory access and the attempted printing of non string data.

5.2 Limitations and future work

5.2.1 Permissions

In order to explore possible problems nasty programs were constructed. These have been used during testing to improve the loader. There is however another class of programs, malicious programs which attempt to access memory which was allocated for another loaded program or even the loader itself. Due to the lack of memory protection methods beyond the normal read write and execute permissions all programs share these permissions. As such any program could take control of most other programs.

A means of protecting loaded programs from each other is documented in [3]. The proposed protection system would enable fine grain access control and secure the loader from ill written programs and malicious

5.3 Applications

The loader offers a platform which could be extended to allow dynamic task execution and placement. A shell program could be used to offer a dynamic interface. Combined with other programs and daemons a simple operating system could be realized.

Conclusions

6.1 Benchmarking results

In benchmarks we have seen that performance... It quite likely works

6.2 Stability

The loader lacks some of the protection mechanisms required for the stable execution of untrusted code. However under normal execution the loader is quite stable and handles any errors it can by terminating the offending program and offering several handles for debugging purposes.

During testing ??? programs have been run in tandem. ??? programs have been run on the same core. ??? programs have been run in a recursive manner.

6.3 Final conclusion

The system...

6.4 References

Bibliography

- [1] Tool Interface Standard (TIS) Committee. *Executable and Linking Format (ELF) Specification, Version 1.2*, May 1995.
- [2] John D. Polstra. *FreeBsd Source src/sys/sys/elf64.h, version 1.9, 5 May 2012*, <http://freebsd.active-venture.com/FreeBSD-src/tree/news/src/sys/elf64.h.html>, September 1999.
- [3] Emmett Witchel, Josh Cates, and Krste Asanović. Mondrian memory protection. In *Proc. 10th international conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS'02)*, ASPLOS-X, pages 304–316, New York, NY, USA, 2002. ACM.

Problems and in depth solutions

A.1 Bugs

A.1.1 String data relocation

At an early stage in the loaders progress data relocation was done at compile time. The assumption was made the compiler would generate code to correct data pointers included in initialized variables. However this is not the case as was concluded when a simple program designed to print an array of strings was run and it became clear that these strings were assumed to be at a fixed location. The error in the location could be expressed as the loaded programs base. The programs continued to show this behavior when compiled with the -fPIC and -fpic compiler flags.

In order to solve this problem, which is a symptom of an incomplete relocation process the data pointers need to be corrected. In order to know which data needs to be corrected for the actual base the compiler needs to be told that our loader will finish the loading process. This is done by passing it the -shared flag. This flag prevents the compiler from incorrectly assuming a value for the base address and include relocation information into the produced ELF file.

Example of the printy programs reported sections when not compiled with -shared

Section	#(Type):	Name, Type#	Flags,	Addr,	Off,
Section 0	(NULLTYPE):	, 0,	0,	0,	0,
Section 1	(Progbits):	.text, 1,	6,	16777216,	65536,
Section 2	(Progbits):	.rodata, 1,	2,	16778048,	66368,
Section 3	(Progbits):	.eh_frame_hdr, 1,	2,	16778112,	66432,
Section 4	(Progbits):	.eh_frame, 1,	2,	16778136,	66456,
Section 5	(Progbits):	.got, 1,	3,	16843720,	66504,
Section 6	(Nobits):	.bss, 8,	3,	16843720,	66504,
Section 7	(Progbits):	.comment, 1,	0,	0,	66504,
Section 8	(Progbits):	.argroom, 1,	0,	0,	66522,
Found the argument section					
Section 9	(Strtab):	.shstrtab, 3,	0,	0,	74714,
Section 10	(Symtab):	.symtab, 2,	0,	0,	75576,
Section 11	(Strtab):	.strtab, 3,	0,	0,	75984,

Spawning program from 0x80101230 of size 0x1290e with flags 2
 To cores: 1 @ 0
 Returning from Loader main

The same program With -shared

Section	#(Type):	Name, Type#	Flags,	Addr,	Off,
Section 0	(NULLTYPE):	, 0,	0,	0,	0,
Section 1	(Progbits):	.text, 1,	6,	16777216,	65536,

```

Section 2( Hash):      .hash,      5,      2,      400,      400,
Section 3( Dynsym):    .dynsym,    11,      2,      576,      576,
Section 4( Strtab):    .dynstr,     3,      2,      792,      792,
Section 5( Rela):      .rela.plt,   4,      2,      848,      848,
Section 6(Progbits):   .rodata,     1,      2,    16778048,    66368,
Section 7(Progbits):   .eh_frame_hdr, 1,      2,    16778112,    66432,
Section 8(Progbits):   .eh_frame,   1,      2,    16778136,    66456,
Section 9( Dynamic):   .dynamic,    6,      3,    16843720,    66504,
Section 10(Progbits):  .plt,        1,      7,    16844016,    66800,
Section 11(Progbits):  .got,         1,      3,    16844064,    66848,
Section 12( Nobits):   .bss,         8,      3,    16844072,    66856,
Section 13(Progbits):  .comment,     1,      0,         0,    66856,
Section 14(Progbits):  .argroom,     1,      0,         0,    66874,
Found the argument section
Section 15( Strtab):    .shstrtab,    3,      0,         0,    75066,
Section 16( Symtab):    .symtab,      2,      0,         0,    76352,
Section 17( Strtab):    .strtab,      3,      0,         0,    76976,
Spawning program from 0x80101230 of size 0x12d27 with flags 2
To cores: 1 @ 0
Returning from Loader main

```

As can be observed, the RELA section.

A.2 Example Configurations

These configurations load a single program each with the specified arguments, environment and specific loader settings.

```

verbose=true
filename=/path/to/file/elffile

```

```

argv1
argv2
argv3

```

```

env0=1
env1=cookies
env3=needs more ducktape
env4=sudo su
env5=make sandwich -j9001

```